

```
In [35]: # Nmae : Himanshu Agarwal , Net id: HXA180027 ,
#q2.3 Titanic dataset
#mporting Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```


```
In [36]: #import data using the pandas libraries.
titanic3_dataset = pd.read_csv('titanic3.csv')
titanic3_dataset.head()
```

Out[36]:

	pclass	survived	name	sex	age	sibsp	parch	ticket	fare	cabin	embarked
0	1	1	Allen, Miss. Elisabeth Walton	female	29.00	0	0	24160	211.3375	B5	S
1	1	1	Allison, Master. Hudson Trevor	male	0.92	1	2	113781	151.5500	C22 C26	S
2	1	0	Allison, Miss. Helen Lorraine	female	2.00	1	2	113781	151.5500	C22 C26	S
3	1	0	Allison, Mr. Hudson Joshua Creighton	male	30.00	1	2	113781	151.5500	C22 C26	S
4	1	0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female	25.00	1	2	113781	151.5500	C22 C26	S

```
In [37]: #Heatmap to check the missig values
sns.heatmap(titanic3_dataset.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

Out[37]: <matplotlib.axes._subplots.AxesSubplot at 0x1a45ef8ab08>



In []:

```
In [38]: # create a correlation matrix that measures the linear relationships between the variables.
#The correlation matrix can be formed by using the corr function from the pandas dataframe library.
#We will use the heatmap function from the seaborn library to plot the correlation matrix
correlation_matrix = titanic3_dataset.corr().round(2)
sns.heatmap(data=correlation_matrix, annot=True)
```

Out[38]: <matplotlib.axes._subplots.AxesSubplot at 0x1a45e954188>




In [39]: *#Through Correlation matrix we can see that few columns are not present which shows they are independent and has no correlation.
#We can drop those features as they will not play any role in deciding the target feature.
#So dropping features 'Cabin', 'body', 'boat', and 'home.dest'.Same way name and ticket.*

```
In [40]: titanic3_dataset.drop('cabin',axis=1,inplace=True)
titanic3_dataset.drop('body',axis=1,inplace=True)
titanic3_dataset.drop('boat',axis=1,inplace=True)
titanic3_dataset.drop('home.dest',axis=1,inplace=True)
```

```
In [41]: #Heatmap to check the missig values
sns.heatmap(titanic3_dataset.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```


Out[41]: <matplotlib.axes._subplots.AxesSubplot at 0x1a45ea76fc8>



In [42]: *#Roughly 20 percent of the Age data is missing. The proportion of Age missing is likely small enough for
#reasonable replacement with some form of imputation.*


```
In [43]: #Data Visualization
#using count plot checking how many passangers survived
sns.set_style('whitegrid')
sns.countplot(x='survived',data=titanic3_dataset,palette='RdBu_r')
```

Out[43]: <matplotlib.axes._subplots.AxesSubplot at 0x1a45e91a648>




```
In [44]: #using count plot checking how many passangers of different genders survived
sns.set_style('whitegrid')
sns.countplot(x='survived',hue='sex',data=titanic3_dataset,palette='RdBu_r')
```

Out[44]: <matplotlib.axes._subplots.AxesSubplot at 0x1a45eb1c8c8>



```
In [45]: #Data cleaning
#using box plot checking class wise average age of passangers
plt.figure(figsize=(12, 7))
sns.boxplot(x='pclass',y='age',data=titanic3_dataset,palette='winter')
```

Out[45]: <matplotlib.axes._subplots.AxesSubplot at 0x1a45eb4bc08>



```
In [46]: #fill in missing age data instead of just dropping the missing age data rows
#One way to do this is by filling in the mean age of all the passengers (imputation).
#However we can be smarter about this and check the average age by passenger class.
```

```
In [47]: #Method to get the average age for a class
```

```
def impute_age(cols):
    Age = cols[0]
    Pclass = cols[1]

    if pd.isnull(Age):

        if Pclass == 1:
            return 37

        elif Pclass == 2:
            return 29

        else:
            return 24

    else:
        return Age
```

```
In [48]: #imputing the missing age values for each class.
titanic3_dataset['age'] = titanic3_dataset[['age', 'pclass']].apply(impute_age,
axis=1)

#to see the count of the missing values left in eachh column
titanic3_dataset.isnull().sum()
```

```
Out[48]: pclass      0
survived    0
name        0
sex         0
age         0
sibsp       0
parch       0
ticket      0
fare        1
embarked    2
dtype: int64
```

```
In [49]: #As the number of rows with null data for fare and embarked are too low as com
pared to total size of the dataset, so we can drop them.
titanic3_dataset.dropna(inplace=True)
```

```
In [50]: #We'll need to convert categorical features to dummy variables using pandas! L
ike 'sex' and 'embarked' features.
#Otherwise our machine learning algorithm won't be able to directly take in th
ose features as inputs.
#the 'Name' and 'Ticket' column have no relationship with whether the person s
urvived or not,observed using correlation matrix.
#So we drop these 2 columns and we convert the other two columns into numerica
l values
```

```
In [51]: sex = pd.get_dummies(titanic3_dataset['sex'],drop_first=True)
embark = pd.get_dummies(titanic3_dataset['embarked'],drop_first=True)
titanic3_dataset.drop(['sex','embarked','name','ticket'],axis=1,inplace=True)
```

```
In [52]: #Data is ready to apply Logistic regression
titanic3_dataset.head()
```

```
Out[52]:
```

	pclass	survived	age	sibsp	parch	fare
0	1	1	29.00	0	0	211.3375
1	1	1	0.92	1	2	151.5500
2	1	0	2.00	1	2	151.5500
3	1	0	30.00	1	2	151.5500
4	1	0	25.00	1	2	151.5500

```
In [53]: #Preprocessing of Data for Logistic regression
```

```
In [54]: #train the model with 80% of the samples and test with the remaining 20%.
#We do this to assess the model's performance on unseen data.
n = int(len(titanic3_dataset)*0.80)
#splitting data in training and test set
df_train, df_test = titanic3_dataset.iloc[:n, :], titanic3_dataset.iloc[n:, :]
titanic3_dataset.columns
```

```
Out[54]: Index(['pclass', 'survived', 'age', 'sibsp', 'parch', 'fare'], dtype='object')
```

```
In [55]: # Initial Coefficients
B = np.array([0, 0, 0, 0, 0, 0]) #Weights array
alpha = 0.0001 # Learning rate
```

```
In [ ]:
```

```
In [56]: #Splitting the training and testing data in X,Y train and test sets.
dt = df_train
m = len(df_train.iloc[:, :-1])
x0 = np.ones(m)
Xtrain = np.array([x0, dt['pclass'], dt['age'], dt['sibsp'], dt['parch'], dt['fare']]).T
ytrain = np.array(dt['survived'])

m = len(df_test.iloc[:, :-1])
x0 = np.ones(m)
dt = df_test
Xtest = np.array([x0, dt['pclass'], dt['age'], dt['sibsp'], dt['parch'], dt['fare']]).T
ytest = np.array(dt['survived'])
```

```
In [57]: #Cost function
def cost_function(X, Y, B):
    m = len(Y)
    J = np.sum((X.dot(B) - Y) ** 2)/(2 * m)
    return J
```

```
In [58]: #Sigmoid fuction to get the result in 0-1 range
def sigmoid(z):
    return 1 / (1 + np.exp(-z))
```

```

In [59]: #mplementation of SGD Logistic regression with momentum for the given numbers
         of iterations.
def gradient_descent(titanic3_data, B, alpha, iterations):
    cost_history = [0] * iterations
    vt=0
    k=10
    for iteration in range(iterations):
        # Sampling the dataset, getting k random records
        dt= titanic3_data.sample(k)
        dt.head()
        m = k
        x0 = np.ones(m)
        #Getting X and Y, from the random sample of size k.
        X1 = np.array([x0, dt['pclass'],dt['age'],dt['sibsp'],dt['parch'],dt[
'fare']] ).T
        Y1 = np.array(dt['survived'])

        z = X1.dot(B)
        h = sigmoid(z) #Calculating the sigmoid of the hypothesis
        # Difference b/w Hypothesis and Actual Y
        loss = h - Y1
        # Gradient Calculation
        gamma = 0.9
        vt = (gamma*vt) + alpha * (X1.T.dot(loss)/k) #adding momentum
        # Changing Values of B using Gradient
        B = B - vt
        # New Cost Value
        cost = cost_function(X1, Y1, B)
        cost_history[iteration] = cost

    return B, cost_history

```

```

In [60]: alpha= 0.001 #Learning rate
         iterations=10000
         theta , cost = gradient_descent(df_train, B, alpha, iterations)

```

```

In [61]: #Method to predict the value of output feature based on the weights received a
         fter training.
def predict(X, theta):
    return sigmoid(np.dot(X, theta)).round()

```

```

In [ ]:

```

```

In [62]: #using scikit-learn's Libraries to get the reports.
         from sklearn.metrics import classification_report
         from sklearn.metrics import accuracy_score

```



```
In [63]: predictions_train = predict(Xtrain,theta) # Making prediction on training set
print("Classification Report on Training data")
print(classification_report(ytrain,predictions_train))
print("Accuracy on Training data:",accuracy_score(ytrain, predictions_train))
```

```
Classification Report on Training data
              precision    recall  f1-score   support

     0       0.58         1.00         0.73         601
     1       0.00         0.00         0.00         443

 accuracy          0.58         0.58         0.58         1044
 macro avg         0.29         0.50         0.37         1044
 weighted avg         0.33         0.58         0.42         1044
```

Accuracy on Training data: 0.5756704980842912

C:\Users\himan\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1437: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.
'precision', 'predicted', average, warn_for)

```
In [64]: predictions_test = predict(Xtest,theta) # Making prediction on test set
print("Classification Report on test data")
print(classification_report(ytest,predictions_test))
print("Accuracy on test data:",accuracy_score(ytest, predictions_test))
```

```
Classification Report on test data
              precision    recall  f1-score   support

     0       0.79         1.00         0.88         207
     1       0.00         0.00         0.00          55

 accuracy          0.79         0.79         0.79         262
 macro avg         0.40         0.50         0.44         262
 weighted avg         0.62         0.79         0.70         262
```

Accuracy on test data: 0.7900763358778626

In []:

In []:

In []:

In []:

In []:

In []: